

DTS202TC Foundation of Parallel Computing

Lecture 3

Instructor: Md Maruf Hasan

Module Leader: Hong-Bin Liu

AIAC, Xi'an Jiaotong-Liverpool University

November 2023

- Assessment groups have been released on the Learning Mall
- Check and Confirm before the deadline.

Goals for this week



Xi'an Jiaotong-Liverpool University
西交利物浦大学

- Profiling tools (Review)
- Parallel Programming with **Pthreads**.

-
- Profiling Tools
 - Pthreads
 - Background
 - Pthreads Basics
 - Our First Pthreads Program
 - Fundamental Concepts and Common Functions
 - Matrix-Vector Multiplication
 - Critical Sections and Synchronization
 - Critical Sections
 - Busy-Waiting
 - Mutex
 - Semaphore
 - Barrier and Condition Variables
 - Thread Safety

Profiling tools (profilers)

Wikipedia: Profiling is a form of dynamic program analysis that measures, for example, the space (memory) or time complexity of a program, the usage of particular instructions, or the frequency and duration of function calls. Most commonly, profiling information serves to aid program optimization, and more specifically, performance engineering.

- Gprof, Xcode Instruments
- mpiP
- HPCToolkit, caliper
- Even, MATLAB got a Profiler too!

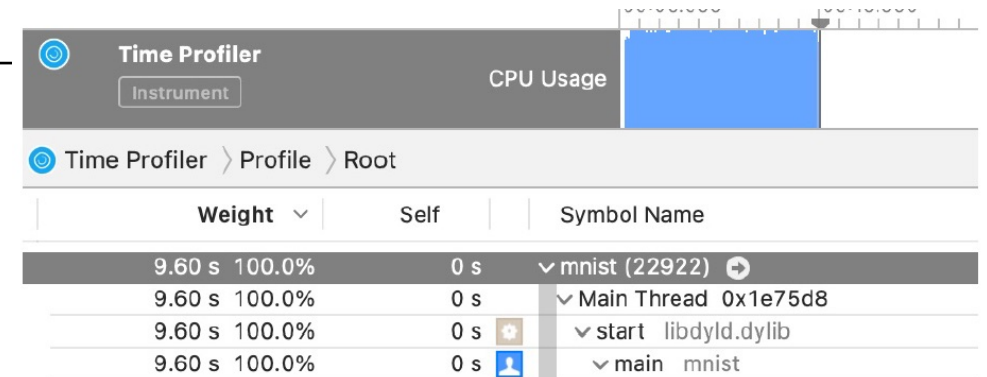


Figure 1: Instruments on Mac

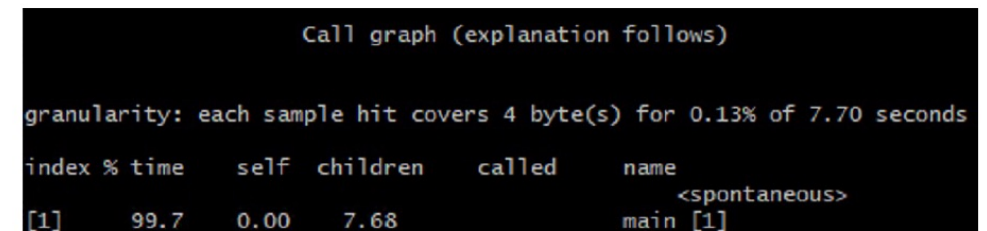


Figure 2: gprof on Cygwin Windows

GNU gprof Tutorial:

<https://www.thegeekstuff.com/2012/08/gprof-tutorial/>

How to Profile a C program in Linux using GNU gprof

<https://www.maketecheasier.com/profile-c-program-linux-using-gprof/>

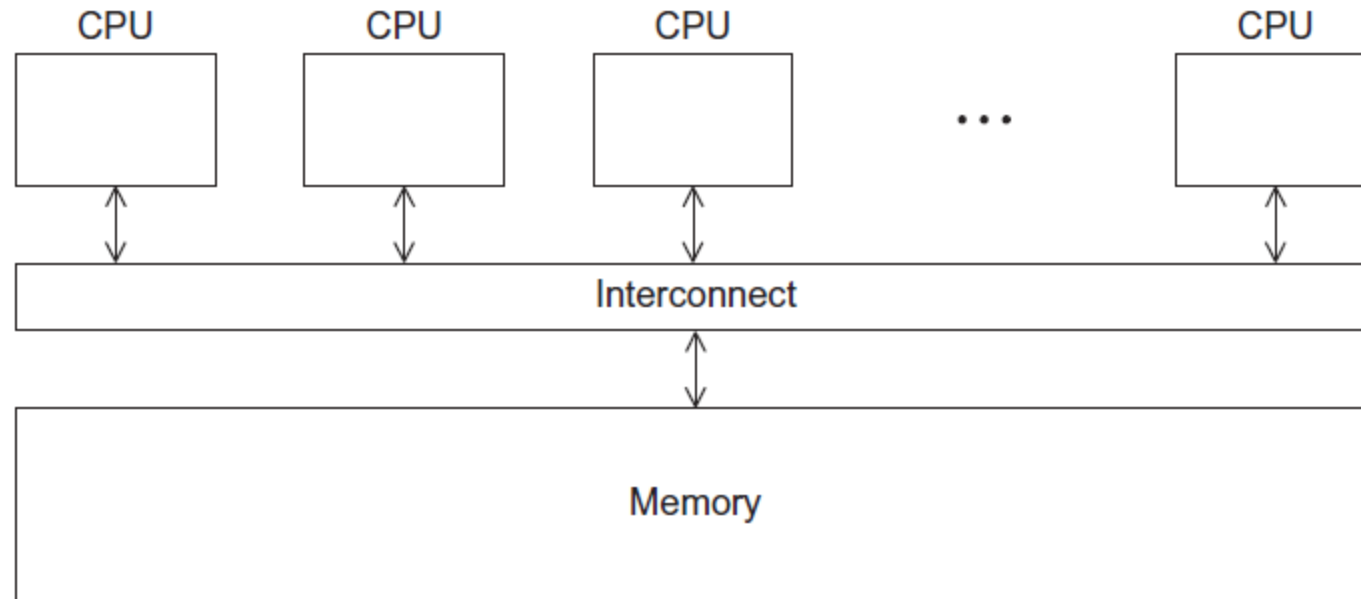
Covered in the Tutorial Session. Important Stuff. Why?

Performance Evaluation, Justification

A profiler is a **Microscope** for a Programmer!

-
- Background
 - Pthreads Basics
 - Our First Pthreads Program
 - Fundamental Concepts and Common Functions
 - Matrix-Vector Multiplication
 - Critical Sections and Synchronization
 - Critical Sections
 - Busy-Waiting
 - Mutex
 - Semaphore
 - Barrier and Condition Variables
 - Thread Safety

A Shared Memory System



Threads and Processes

- In shared-memory programming, a process is an instance of a running (or suspended) program.
- Threads are analogous to a “light-weight” process.
- In a shared-memory program, a single process may have multiple threads of control.

- IEEE had a **POSIX** 1003 group that defined an interface to multithreaded programming
 - This is called Pthreads (POSIX Thread). **Portable**
 - A standard for Unix-like operating systems. The Pthreads API is only available on POSIX systems, however, we are using Cygwin on Windows or Unix Virtual Machine.
 - A library that can be linked with C programs.
 - Provides **primitives for thread management and synchronization**
- Threads are peers, no parent/child relationship

-
- Background
 - Pthreads Basics
 - Our First Pthreads Program
 - Fundamental Concepts and Common Functions
 - Matrix-Vector Multiplication
 - Critical Sections and Synchronization
 - Critical Sections
 - Busy-Waiting
 - Mutex
 - Semaphore
 - Barrier and Condition Variables
 - Thread Safety

Hello World!



```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

/* Global variable: accessible to all threads */
int thread_count;

void *Hello(void* rank); /* Thread function */

int main(int argc, char* argv[]) {
    long thread; /* Use long in case of a 64-bit system */
    pthread_t* thread_handles;

    /* Get number of threads from command line */
    thread_count = strtol(argv[1], NULL, 10);

    thread_handles = malloc (thread_count*sizeof(pthread_t));
```

declares the various Pthreads functions, constants, types, etc.

global variable

Hello World! Cont.



```
for (thread = 0; thread < thread_count; thread++)  
    pthread_create(&thread_handles[thread], NULL,  
                  Hello, (void*) thread);
```

Create and start new threads

```
printf("Hello from the main thread\n");
```

```
for (thread = 0; thread < thread_count; thread++)  
    pthread_join(thread_handles[thread], NULL);
```

Join threads

```
free(thread_handles);  
return 0;  
} /* main */
```


```
void *Hello(void* rank) {  
    long my_rank = (long) rank; /* Use long in case of 64-bit system */  
  
    printf("Hello from thread %ld of %d\n", my_rank, thread_count);  
  
    return NULL;  
} /* Hello */
```

Compiling a Pthread program



```
gcc -g -Wall -o pthread_hello pthread_hello.c -lpthread
```

link in the Pthreads library



Running a Pthreads program



```
./pth_hello <number of threads>
```

```
./pth_hello 1
```

```
Hello from the main thread  
Hello from thread 0 of 1
```

```
./pth_hello 4
```

```
Hello from the main thread  
Hello from thread 0 of 4  
Hello from thread 1 of 4  
Hello from thread 2 of 4  
Hello from thread 3 of 4
```

- thread_count
- Can introduce subtle and confusing bugs!
- Limit use of global variables to situations in which they're really needed.
 - Shared variables.


-
- Pthreads Basics
 - Our First Pthreads Program
 - Fundamental Concepts and Common Functions
 - Matrix-Vector Multiplication
 - Critical Sections and Synchronization
 - Critical Sections
 - Busy-Waiting
 - Mutex
 - Semaphore
 - Barrier and Condition Variables
 - Read-Write Locks
 - Thread Safety

- In Pthreads, the threads are started by the program executable.
 - Include code to explicitly start the threads
 - Need data structures to store threads information
- Processes in MPI are usually started by a script.

Starting the Threads




pthread.h



pthread_t

**One object for
each thread.**



```
int pthread_create (  
    pthread_t*  thread_p /* out */,  
    const pthread_attr_t* attr_p /* in */,  
    void* (*start_routine) ( void ) /* in */,  
    void* arg_p /* in */ );
```

The function creates and starts a new thread.

- **Opaque**
- The actual data that they store is system-specific.
- Their data members aren't directly accessible to user code.
- However, the Pthreads standard guarantees that a pthread_t object does store enough information to uniquely identify the thread with which it's associated.

A Closer Look (1)



```
int pthread_create (  
    pthread_t* thread_p /* out */,  
    const pthread_attr_t* attr_p /* in */,  
    void* (*start_routine) ( void ) /* in */,  
    void* arg_p /* in */ );
```

We won't be using, so we just pass NULL.

Allocate before calling.

A Closer Look (2)



```
int pthread_create (  
    pthread_t* thread_p /* out */,  
    const pthread_attr_t* attr_p /* in */,  
    void* (*start_routine) ( void ) /* in */,  
    void* arg_p /* in */ );
```

Pointer to the argument that should
be passed to the function *start_routine*.

The function that the thread is to run.

Function started by pthread_create

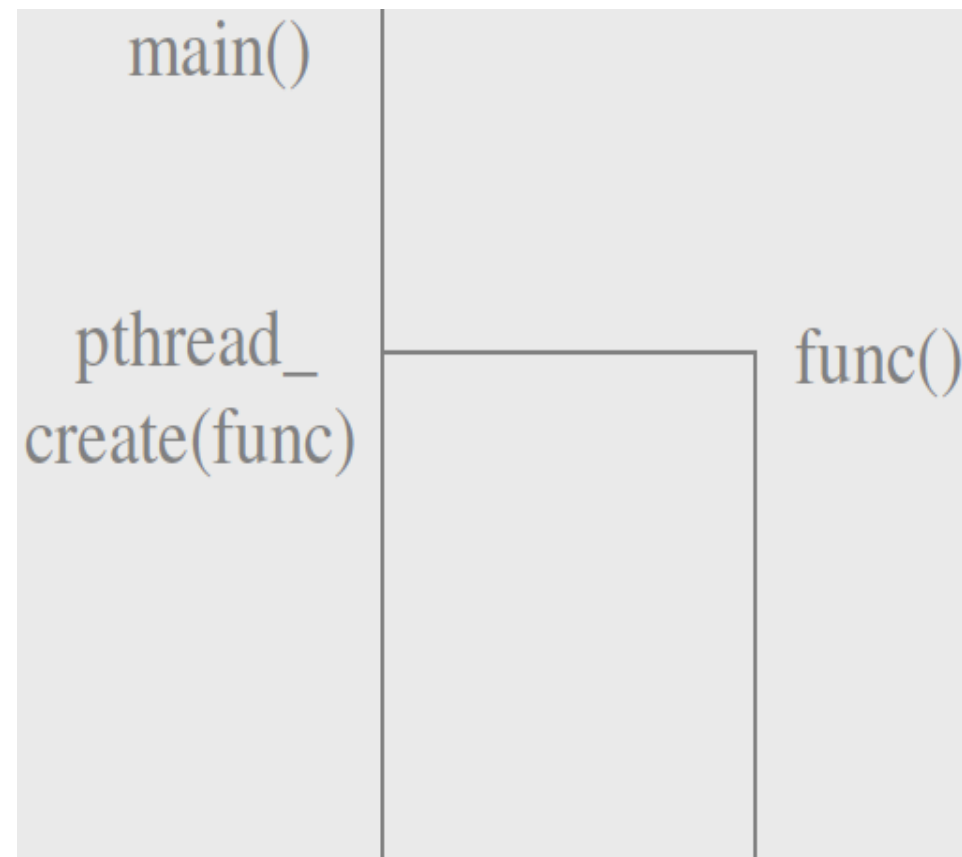


- Prototype:
`void* thread_function (void* args_p) ;`
- Void* can be cast to any pointer type in C.
- So args_p can point to a list containing one or more values needed by thread_function.
- Similarly, the return value of thread_function can point to a list of one or more values.

Example of Thread Creation



```
void *func(void *arg) {  
    int *i=arg;  
    .....  
}  
void main()  
{  
    int X; pthread_t id;  
    ....  
    pthread_create(&id, NULL,  
                  func, &X);  
    ...  
}
```



- We call the function `pthread_join` once for each thread.
- ```
int pthread_join(pthread_t new_id, /* in */
 void ** ret_val_p /* out */)

```

  - Waits for the thread with identifier `new_id` to terminate, either by returning or by calling `pthread_exit()`
- A single call to `pthread_join` will wait for the thread associated with the `pthread_t` object to complete.


# Thread Joining Example

---




```
void *func(void *) { }
pthread_t id; int X;
pthread_create(&id, NULL, func, &X);
.....
pthread_join(id, NULL);
.....
```

# Running the Threads



Main thread forks and joins two threads.

# Pthreads Programming Model




```
pthread_t pthread_self(void);
```

- To determine the thread ID of the calling thread

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - **Matrix-Vector Multiplication**
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
  - Thread Safety

# Matrix-vector Multiplication



**FIGURE 3.11**

Matrix-vector multiplication

$$y_i = \sum_{j=0}^{n-1} a_{ij}x_j$$


```
/* For each row of A */
for (i = 0; i < m; i++) {
 y[i] = 0.0;
 /* For each element of the row and each element of x */
 for (j = 0; j < n; j++)
 y[i] += A[i][j]* x[j];
}
```



# Using Three Pthreads



| Thread | Components<br>of y |
|--------|--------------------|
| 0      | y[0], y[1]         |
| 1      | y[2], y[3]         |
| 2      | y[4], y[5]         |



thread 0

```
y[0] = 0.0;
for (j = 0; j < n; j++)
 y[0] += A[0][j]*x[j];
```



general case

```
y[i] = 0.0;
for (j = 0; j < n; j++)
 y[i] += A[i][j]*x[j];
```

# Pthreads Matrix-vector Multiplication



```
void *Pth_mat_vect(void* rank) {
 long my_rank = (long) rank;
 int i, j;
 int local_m = m/thread_count;
 int my_first_row = my_rank*local_m;
 int my_last_row = (my_rank+1)*local_m - 1;

 for (i = my_first_row; i <= my_last_row; i++) {
 y[i] = 0.0;
 for (j = 0; j < n; j++)
 y[i] += A[i][j]*x[j];
 }

 return NULL;
} /* Pth_mat_vect */
```

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
  - Thread Safety

# Data Race in Pthreads Program



```
static int s = 0;
```

Thread 1

```
for i = 0, n/2-1
 s = s + f(A[i])
```

Thread 2

```
for i = n/2, n-1
 s = s + f(A[i])
```

- Problem is a **race condition** on variable **s** in the program
- When multiple threads attempt to access a shared resource, at least one of the accesses is an update, and the accesses can result in an error, we have a **race condition**.

$$\pi = 4 \left( 1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \cdots + (-1)^n \frac{1}{2n+1} + \cdots \right)$$

```
double factor = 1.0;
double sum = 0.0;
for (i = 0; i < n; i++, factor = -factor) {
 sum += factor/(2*i+1);
}
pi = 4.0*sum;
```

# A thread function for computing $\pi$



```
void* Thread_sum(void* rank) {
 long my_rank = (long) rank;
 double factor;
 long long i;
 long long my_n = n/thread_count;
 long long my_first_i = my_n*my_rank;
 long long my_last_i = my_first_i + my_n;

 if (my_first_i % 2 == 0) /* my_first_i is even */
 factor = 1.0;
 else /* my_first_i is odd */
 factor = -1.0;

 for (i = my_first_i; i < my_last_i; i++, factor = -factor) {
 sum += factor/(2*i+1);
 }

 return NULL;
} /* Thread_sum */
```

# Using a dual core processor

|           | $n$     |          |           |            |
|-----------|---------|----------|-----------|------------|
|           | $10^5$  | $10^6$   | $10^7$    | $10^8$     |
| $\pi$     | 3.14159 | 3.141593 | 3.1415927 | 3.14159265 |
| 1 Thread  | 3.14158 | 3.141592 | 3.1415926 | 3.14159264 |
| 2 Threads | 3.14158 | 3.141480 | 3.1413692 | 3.14164686 |

- Note that as we increase  $n$ , the estimate with one thread gets better and better.
- However, for larger values of  $n$ , the result computed by two threads actually gets worse.
- It matters if multiple threads try to update a single shared variable.

# Possible race condition



```
y = Compute(my rank);
x = x + y;
y: private, x: shared, initialized to 0
```

| Time | Thread 0                       | Thread 1                       |
|------|--------------------------------|--------------------------------|
| 1    | Started by main thread         |                                |
| 2    | Call Compute ()                | Started by main thread         |
| 3    | Assign y = 1                   | Call Compute ()                |
| 4    | Put x=0 and y=1 into registers | Assign y = 2                   |
| 5    | Add 0 and 1                    | Put x=0 and y=2 into registers |
| 6    | Store 1 in memory location x   | Add 0 and 2                    |
| 7    |                                | Store 2 in memory location x   |

A race condition or data race occurs when:

- Two processors (or two threads) access the same variable, and at least one does a write.
- The accesses are concurrent (not synchronized) so they could happen simultaneously.



- **Critical Section** is a block of code that updates a shared resources that can only be updated by one thread a time.

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
    - Read-Write Locks
  - Thread Safety

- Create/exit/join
  - provide some form of synchronization,
  - at a very coarse level,
  - requires thread creation/destruction.
- Need for finer-grain synchronization
  - mutex locks
  - condition variables
  - ....
- PTHREADS provides a variety of synchronization facilities for threads to cooperate in accessing shared resources.

```
y = Compute(my_rank);
while (flag != my_rank);
x = x + y;
flag++;
```

flag initialized to 0 by main thread

- Busy-waiting is not an ideal solution to the problem of controlling access to a critical section.
- Since thread 1 will execute the test over and over until thread 0 executes flag++ , if thread 0 is delayed, thread 1 will simply “spin” on the test, eating up CPU cycles. This can be positively disastrous for performance.

- A thread repeatedly tests a condition, but, effectively, does no useful work until the condition has the appropriate value.
- The busy-wait solution would work “provided the statements are executed exactly as they’re written.”
- Beware of optimizing compilers, though!

```
y = Compute(my_rank);
while (flag != my_rank);
x = x + y;
flag++;
```

```
y = Compute(my_rank);
x = x + y;
while (flag != my_rank);
flag++;
```

- Turning off compiler optimizations can seriously degrade performance.

```
void* Thread_sum(void* rank) {
 long my_rank = (long) rank;
 double factor;
 long long i;
 long long my_n = n/thread_count;
 long long my_first_i = my_n*my_rank;
 long long my_last_i = my_first_i + my_n;

 if (my_first_i % 2 == 0)
 factor = 1.0;
 else
 factor = -1.0;

 for (i = my_first_i; i < my_last_i; i++, factor = -factor) {
 while (flag != my_rank);
 sum += factor/(2*i+1);
 flag = (flag+1) % thread_count;
 }

 return NULL;
} /* Thread_sum */
```

When  $n=10^8$ , the serial sum is consistently faster than the parallel sum.

```
void* Thread_sum(void* rank) {
 long my_rank = (long) rank;
 double factor, my_sum = 0.0;
 long long i;
 long long my_n = n/thread_count;
 long long my_first_i = my_n*my_rank;
 long long my_last_i = my_first_i + my_n;

 if (my_first_i % 2 == 0)
 factor = 1.0;
 else
 factor = -1.0;

 for (i = my_first_i; i < my_last_i; i++, factor = -factor)
 my_sum += factor/(2*i+1);

 while (flag != my_rank);
 sum += my_sum;
 flag = (flag+1) % thread_count;

 return NULL;
} /* Thread_sum */
```

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - **Mutex**
    - Semaphore
    - Barrier and Condition Variables
    - Read-Write Locks
  - Thread Safety



- A thread that is busy-waiting may continually use the CPU accomplishing nothing.
- Mutex (mutual exclusion) is a special type of **variable** that can be used to restrict access to a critical section to a single thread at a time.
- Used to guarantee that one thread “excludes” all other threads while it executes the critical section.

- The Pthreads standard includes a special type for mutexes: `pthread_mutex_t`. Need to be initialized by the system before it's used.

```
int pthread_mutex_init(
 pthread_mutex_t* mutex_p /* out */
 const pthread_mutexattr_t* attr_p /* in */);
```

- Creates a new mutex lock.
- Attribute: normal, recursive, errorcheck  
(Ignore the second argument in this book, just pass NULL.)

- When a Pthreads program finishes using a mutex, it should call `pthread_mutex_destroy`

```
int pthread_mutex_destroy(pthread_mutex_t* mutex_p /* in/out */);
```

- In order to gain access to a critical section a thread calls `pthread_mutex_lock`

```
int pthread_mutex_lock(pthread_mutex_t* mutex_p /* in/out */);
```

- Tries to acquire the lock specified by mutex.
- If mutex is already locked, then calling thread blocks until mutex is unlocked.

- When a thread is finished executing the code in a critical section, it should call `pthread_mutex_unlock`

```
int pthread_mutex_unlock(pthread_mutex_t* mutex_p /* in/out */);
```

- If calling thread has mutex currently locked, this will unlock the mutex.
- If other threads are blocked waiting on this mutex, one will unblock and acquire mutex, which one is determined by the scheduler.

# Global Sum Function that Uses a Mutex



```
void* Thread_sum(void* rank) {
 long my_rank = (long) rank;
 double factor;
 long long i;
 long long my_n = n/thread_count;
 long long my_first_i = my_n*my_rank;
 long long my_last_i = my_first_i + my_n;
 double my_sum = 0.0;

 if (my_first_i % 2 == 0)
 factor = 1.0;
 else
 factor = -1.0;

 for (i = my_first_i; i < my_last_i; i++, factor = -factor) {
 my_sum += factor/(2*i+1);
 }

 pthread_mutex_lock(&mutex);
 sum += my_sum;
 pthread_mutex_unlock(&mutex);

 return NULL;
} /* Thread_sum */
```

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
    - Read-Write Locks
  - Thread Safety

- 
- Busy-waiting enforces the order threads access a critical section.
  - Using mutexes, the order is left to chance and the system.
  - There are applications where we need to control the order threads access the critical section.

# Problems with a Mutex Solution



```
/* n and product_matrix are shared and initialized by the main
thread */
/* product_matrix is initialized to be the identity matrix */
```

```
void* Thread_work(void* rank) {
 long my_rank = (long) rank;
 matrix_t my_mat = Allocate_matrix(n);
 Generate_matrix(my_mat);
 pthread_mutex_lock(&mutex);
 Multiply_matrix(product_mat, my_mat);
 pthread_mutex_unlock(&mutex);
 Free_matrix(&my_mat);
 return NULL;
} /* Thread_work */
```



# A First Attempt at Sending Messages Using Pthreads



```
/* messages has type char**. It's allocated in main. */
/* Each entry is set to NULL in main. */
void *Send_msg(void* rank) {
 long my_rank = (long) rank;
 long dest = (my_rank + 1) % thread_count;
 long source = (my_rank + thread_count - 1) % thread_count;
 char* my_msg = malloc(MSG_MAX*sizeof(char));

 sprintf(my_msg, "Hello to %ld from %ld", dest, my_rank);
 messages[dest] = my_msg;

 if (messages[my_rank] != NULL) {
 printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
 } else {
 printf("Thread %ld > No message from %ld\n", my_rank, source);
 }

 return NULL;
} /* Send_msg */
```

No receiving message for  
multithreads

# Possible Solutions with Problems



**S1**: Change **if** statement to **while** statement:

```
-if (messages[my_rank] != NULL)
 printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
-while (messages[my_rank] = NULL);
 printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
•However, there is a busy-waiting problem.
```

•**S2**: Mutex Lock:

```
...
Pthread_mutex_lock(mutex[dest]);
...
Messages[dest]=msg;
Pthread_mutex_unlock(mutex[dest]);
...
Pthread_mutex_lock(mutex[myrank]);
printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
...
```

•However, it still will crash.

- Semaphores can be thought of as a special type of `unsigned int` , so they can take on the values 0, 1, 2, . . . . When they take on the values 0 and 1, called a binary semaphore.
- For `binary semaphore`, 0 corresponds to a locked mutex, and 1 corresponds to an unlocked mutex.
  - Initialized to 1(unlocked).
  - Before the critical, call `sem_wait` . Executing `sem_wait` will block if the semaphore is 0. If the semaphore is nonzero, it will decrement the semaphore and proceed.
  - After executing the critical, call `sem_post` , which increments the semaphore, and a thread waiting in `sem_wait` can proceed.

# Syntax of the Various Semaphore Functions



```
#include <semaphore.h>
```

Semaphores are not part of  
Pthreads; you need to add this.

```
int sem_init(
 sem_t* semaphore_p /* out */,
 int shared /* in */,
 unsigned initial_val /* in */);
```

Initialize the semaphore descriptor.

Unlock a semaphore.

```
int sem_destroy(sem_t* semaphore_p /* in/out */);
int sem_post(sem_t* semaphore_p /* in/out */);
int sem_wait(sem_t* semaphore_p /* in/out */);
```

Destroy a semaphore.

Lock a semaphore.

# Send Messages Using Semaphores



```
1 /* messages is allocated and initialized to NULL in main */
2 /* semaphores is allocated and initialized to 0 (locked) in
 main */
3 void* Send_msg(void* rank) {
4 long my_rank = (long) rank;
5 long dest = (my_rank + 1) % thread_count;
6 char* my_msg = malloc(MSG_MAX*sizeof(char));
7
8 sprintf(my_msg, "Hello to %ld from %ld", dest, my_rank);
9 messages[dest] = my_msg;
10 sem_post(&semaphores[dest])
11 /* "Unlock" the semaphore of dest */
12
13 /* Wait for our semaphore to be unlocked */
14 sem_wait(&semaphores[my_rank]);
15 printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
16 return NULL;
17 } /* Send_msg */
```

- The message-sending problem didn't involve a critical section.
- Thread my\_rank couldn't proceed until thread source had finished creating the message.
- When a thread can't proceed until another thread has taken some action, is sometimes called **producer-consumer synchronization**.

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
  - Thread Safety

- 
- Synchronizing the threads to make sure that they all are at the same point in a program is called a **barrier**.
  - No thread can cross the barrier until all the threads have reached it.

# Using Barriers to Time the Slowest Thread



```
/* Shared */
double elapsed_time;
. . .
/* Private */
double my_start, my_finish, my_elapsed;
. . .
Synchronize threads;
Store current time in my_start;
/* Execute timed code */
. . .
Store current time in my_finish;
my_elapsed = my_finish - my_start;

elapsed = Maximum of my_elapsed values;
```



# Using Barriers for Debugging

---



```
point in program we want to reach;
barrier;
if (my_rank == 0) {
 printf("All threads reached this point\n");
 fflush(stdout);
}
```

- Implementing a barrier using busy-waiting and a mutex is straightforward.
- We use a shared **counter** protected by the **mutex**.
- When the counter indicates that every thread has entered the critical section, threads can leave the **busy-wait** loop.

# Busy-waiting and a Mutex



```
/* Shared and initialized by the main thread */
```

```
int counter; /* Initialize to 0 */
```

```
int thread_count;
```

```
pthread_mutex_t barrier_mutex;
```

```
. . .
```

```
void* Thread_work(. . .) {
```

```
. . .
```

```
/* Barrier */
```

```
pthread_mutex_lock(&barrier_mutex);
```

```
counter++;
```

```
pthread_mutex_unlock(&barrier_mutex);
```

```
while (counter < thread_count);
```

```
. . .
```

```
}
```

We need one counter variable for each instance of the barrier, otherwise problems are likely to occur.

# Implementing a Barrier with Semaphores



```
/* Shared variables */
int counter; /* Initialize to 0 */
sem_t count_sem; /* Initialize to 1 */
sem_t barrier_sem; /* Initialize to 0 */
. . .
void* Thread_work(...) {
 . . .
 /* Barrier */
 sem_wait(&count_sem);
 if (counter == thread_count - 1) {
 counter = 0;
 sem_post(&count_sem);
 for (j = 0; j < thread_count - 1; j++)
 sem_post(&barrier_sem);
 } else {
 counter++;
 sem_post(&count_sem);
 sem_wait(&barrier_sem);
 }
 . . .
}
```

protects the counter

block threads that have entered the barrier

Reusing  
**barrier\_sem**  
results in a  
race condition.

- A **condition variable** is a data object that allows a thread to suspend execution until a certain event or condition occurs.
- When the event or condition occurs, another thread can **signal** the thread to “wake up.”
- A **condition variable** is always associated with a **mutex**.

```
lock mutex;
if condition has occurred
 signal thread(s);
else {
 unlock the mutex and block;
 /* when thread is unblocked, mutex is relocked */
}
unlock mutex;
```

- Condition variables in Pthreads have type `pthread_cond_t`.

```
int pthread_cond_init(
 pthread_cond_t* cond_p /* out */,
 const pthread_condattr_t* cond_attr_p /* in */);
```

- Creates a new condition variable `cond`.
- Attribute: ignore for now.

```
int pthread_cond_destroy(pthread_cond_t* cond_p /* in/out */);
```

- Destroys the condition variable `cond`.

```
int pthread_cond_signal(pthread_cond_t* cond_var_p /* in/out */);
```

- Unblock one of the blocked threads waiting on cond.
- Which one is determined by scheduler.
- If no thread waiting, then signal is a no-op.


```
int pthread_cond_broadcast(pthread_cond_t* cond_var_p /* in/out */);
```

- Unblock all of the blocked threads waiting on cond.
- If no thread waiting, then signal is a no-op.



```
int pthread_cond_wait(
 pthread_cond_t* cond_var_p /* in/out */,
 pthread_mutex_t* mutex_p /* in/out */);
```

- Blocks the calling thread, waiting on cond.
- Unlocks the mutex.



```
pthread_mutex_unlock(&mutex_p);
wait_on_signal(&cond_var_p);
pthread_mutex_lock(&mutex_p);
```

```
/* Shared */
int counter = 0;
pthread_mutex_t mutex;
pthread_cond_t cond_var;
. . .
void* Thread_work(. . .) {
 . . .
 /* Barrier */
 pthread_mutex_lock(&mutex);
 counter++;
 if (counter == thread_count) {
 counter = 0;
 pthread_cond_broadcast(&cond_var);
 } else {
 while (pthread_cond_wait(&cond_var, &mutex) != 0);
 }
 pthread_mutex_unlock(&mutex);
 . . .
}
```

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
  - Thread Safety

- A block of code is **thread-safe** if it can be simultaneously executed by multiple threads without causing problems.

- 
- Suppose we want to use multiple threads to “tokenize” a file that consists of ordinary English text.
  - The tokens are just contiguous sequences of characters separated from the rest of the text by white-space — a space, a tab, or a newline.

- Divide the input file into lines of text and assign the lines to the threads in a round-robin fashion.
- The first line goes to thread 0, the second goes to thread 1, . . . , the tth goes to thread t, the t + 1st goes to thread 0, etc.
- We can serialize access to the lines of input using semaphores.
- After a thread has read a single line of input, it can tokenize the line using the `strtok` function.

# The Strtok Function



```
char* strtok(
 char* string /* in/out */,
 const char* separators /* in */);
```

- The **idea** is that in the first call, **strtok** caches a pointer to string, and for subsequent calls it returns successive tokens **taken from the cached copy**.
- The first time it's called the string argument should be the text to be tokenized.(One line of input.)
- For subsequent calls, the first argument should be NULL.

# Multi-threaded Tokenizer



```
void* Tokenize(void* rank) {
 long my_rank = (long) rank;
 int count;
 int next = (my_rank + 1) % thread_count;
 char *fg_rv;
 char my_line[MAX];
 char *my_string;

 sem_wait(&sems[my_rank]);
 fg_rv = fgets(my_line, MAX, stdin);
 sem_post(&sems[next]);
 while (fg_rv != NULL) {
 printf("Thread %ld > my line = %s", my_rank, my_line);

 count = 0;
 my_string = strtok(my_line, " \t\n");
 while (my_string != NULL) {
 count++;
 printf("Thread %ld > string %d = %s\n", my_rank, count,
 my_string);
 my_string = strtok(NULL, " \t\n");
 }

 sem_wait(&sems[my_rank]);
 fg_rv = fgets(my_line, MAX, stdin);
 sem_post(&sems[next]);
 }

 return NULL;
} /* Tokenize */
```



- It correctly tokenizes the input stream.

Pease porridge hot.  
Pease porridge cold.  
Pease porridge in the pot  
Nine days old.

# Running with Two Threads



```
Thread 0 > my line = Pease porridge hot.
Thread 0 > string 1 = Pease
Thread 0 > string 2 = porridge
Thread 0 > string 3 = hot.
Thread 1 > my line = Pease porridge cold.
Thread 0 > my line = Pease porridge in the pot
Thread 0 > string 1 = Pease
Thread 0 > string 2 = porridge
Thread 0 > string 3 = in
Thread 0 > string 4 = the
Thread 0 > string 5 = pot
Thread 1 > string 1 = Pease
Thread 1 > my line = Nine days old.
Thread 1 > string 1 = Nine
Thread 1 > string 2 = days
Thread 1 > string 3 = old.
```

Oops!



# What happened?



- `strtok` caches the input line by declaring a variable to have static storage class.
- This causes the value stored in this variable to persist from one call to the next.
- Unfortunately for us, `this cached string is shared, not private`.
- Thus, thread 0's call to `strtok` with the third line of the input has apparently `overwritten` the contents of thread 1's call with the second line.
- So the `strtok` function is not thread-safe. If multiple threads call it simultaneously, the output may not be correct.

# Other Unsafe C Library Functions

---



- Regrettably, it's not uncommon for C library functions to fail to be thread-safe.
- The random number generator `random` in `stdlib.h`.
- The time conversion function `localtime` in `time.h`.

# “re-entrant” (Thread Safe) Functions



- In some cases, the C standard specifies an alternate, thread-safe, version of a function.

```
char* strtok_r(
 char* string /* in/out */,
 const char* separators, /* in */,
 char** saveptr_p /* in/out */);
```

---

Incorrect programs can produce  
correct output!

- Background
- Pthreads Basics
- Matrix-Vector Multiplication
- Critical Sections and Synchronization
  - Critical Sections
  - Busy-Waiting
  - Mutex
  - Semaphore
  - Barrier and Condition Variables
- Thread Safety